

AMJ Retro & JAS Roadmap

Q2 Operating Plan

PART 1

Q1 Retro

We built what was asked — and kept the lights on.

Two primary work areas in AMJ — shipping features & holding the platform together underneath so orders kept flowing.

WHAT WE BUILT

Some of these we could measure, some we couldn't. Most of what we did measure is L0 — top-line movement, not a deep-dive yet. More numbers will come as the analysis catches up.

#	INITIATIVE	WHAT SHIPPED	IMPACT
1	Upfront payment	Payment moved to checkout for non-HA orders.	—
2	Payment Revamp	Rebuilt as a service — init/process/refunds/OMS/recon cons.	12 warehouses, 50% live Jul 9
3	Sub-skip (HA)	Skip orders where substitution uptake / value low.	—
4	Skip Option for NFTC customers	Default opt-in skip of the HA call for NFTC customers; backend logic only.	~150% ↑ daily single skips; 9.2% ↑ total daily skips
5	Voice Bot	Oration AI integrated for FCR-type calls.	Answering 22% of platform call volume
6	Cancellations via CSR	Removed cancel CTA post-warehouse; cancel now via CSR call.	OD per 100 cancel-page views 33.7 → 44.43; incremental GM ₹5.12 L

#	INITIATIVE	WHAT SHIPPED	IMPACT
7	NFS re-enablement (NFS salts)	Sellable again, routed to specialist doctors. Phase 1 = 4 categories.	~₹50 L/month digitised run-rate
8	COD blocking	Stopped showing COD to a high-RTO customer segment.	—
9	Fraud fixes	Found customers gaming the system, plugged the loopholes.	—
10	Multiple DID	Multi-number config + dynamic fallback for Value Meds.	—
11	Prescription fixes	Fixed false "prescription required" errors on valid-Rx orders.	~60% → ~6% of valid orders hitting the error (~10× drop)
12	Doctors Portal — order-flow fix	Released Jun 23. Right CTAs; transfer CTA only during HA call. SOP + training.	—

KEEPING THE LIGHTS ON

Incidents: from **~15 a day** to **~1–2 a day**.

Underneath the build list, this was the quarter's real load. We absorbed **~1,350 incident reports** — ~15 a day — from 255 unique issues, tracing to just **21 root defects**. The same handful of bugs, re-raised over and over, no structured intake, engineers holding the platform together live.

We moved from firefighting live to **structured intake** → **RCA** → **plumbing the root cause**. That's what took ~15 a day down to 1–2.

And this is only what hit the channel — at order level the blast radius is bigger. The worst example: valid orders were wrongly blocked with a "prescription required" error, so the order couldn't be completed. At peak that hit ~1,800 orders a day — ~60% of valid orders (higher still at medicine-SKU level). Now down to **~200 a day, ~6%**.

PART 2

Q2 Roadmap

We start working backwards from outcomes — not forwards from features.

We're not fully there yet. Some work is still foundational, and the customer-experience metric doesn't even exist. So Q2 does two things:

- set up the discovery that gets us fully outcome-led
- anchor what we can to the two L0: **OP2OD** and **CPO**

WHY THESE BETS

We'll be straight: we didn't get into problem discovery deeply in Q1 as we were firefighting. So Q2 isn't a set of proven bets. It is us picking up what we already know matters — what's in flight, what customers keep calling about, and what keeps breaking — and tying each to one of the two outcomes we can measure: OP2OD or CPO.

This quarter leans **heavy on CPO**. The concrete goal: **CSR CPO from ~₹14 to ~₹9 an order**, driven by Voicebot Ph2, Returns & Refunds, Reconnection and Wallet. *(To be vetted and signed off with Finance.)*

On **OP2OD**, **we're building the ammunition**. Between an order being **placed and warehouse-assigned**, we lose orders at several steps — payment, the doctor/HA call, and others — and we can't yet see which loses most. **Discard Intelligence** sizes each drop-off so we attack the biggest, not the loudest. Checkout already goes after the payment loss; Doctor Experience after the call loss.

Improve Order Completion — OP2OD

Checkout & Payment Journey OP2OD · CPO

Why now	Internally cleaned up (new payment service built), but externally we're on an almost-obsolete SDK — Cashfree is barely keeping the lights on. Checkout's scattered: extra clicks, extra page transitions (failure points), no CVV-less, checkout logic still in the frontend.
Problem	We can't see where customers drop. The only number we have is order-level success at 83% — coarse, on already high-intent customers, so it flatters the picture and hides the real step-by-step leak. 55% COD (₹9.17 CH charges/order + ₹3.45 RTO). No payment-step instrumentation.
What will we do	On the new payment service, complete the Juspay integration : pull logic out of the frontend, smoothen the flow, add instrumentation + drop-off visibility.
Juspay unlocks	<i>(vendor claims — to validate)</i> higher SR (routing, auto-retry, downtime fallback, near-zero UPI technical declines); PG-cost reduction; COD-to-prepaid via UPI

QR; PG-agnostic saved cards; 100%-checkout / partial-COD / COD-blocking experiments.

Moves Instrumentation surfaces the real drop-offs (the number may look worse before it looks better — that's the point) → SR levers + fewer failure points optimise **OP2OD**; PG-cost + COD-to-prepaid → **CPO gains**.

CPO reduction

Returns & Refunds CPO

Why now 36 agents run returns and refunds by hand on manual SOPs, handling **~1,500 inbound + ~400 outbound calls** today. A standing cost centre and a known contact load we can attack now.

Problem Every return needs an **agent to approve** — manual SOP, refund decided by hand. Slow, costly, inconsistent, eats CSR bandwidth.

What will we do Self-serve returns: customer raises the return → uploads images → refunded automatically on a rules engine, agent only for exceptions. Big project, delivered in milestones.

Moves Cuts refund-handling cost and agent load → **CPO. Flat-out savings ~₹10–13 L/month**. Faster refunds, fewer calls on the side.

Foundations

Doctor Experience Platform

Why now ~15,000 doctor calls a day — **~54% of our ~28,000 daily orders** touch a doctor. One of the biggest moments in the journey, yet it runs on a portal we treat as pure compliance: high-friction, fragile, the source of most of the incidents above.

Problem We built it as a compliance step, not an experience. Doctors work on a breaking platform; that friction shows up as weaker consult quality for the customer, and as the firefighting load that ate engineering all quarter.

What will we do A product bet from **first principles**: build a platform that makes the doctor's job easy → better consults → better customer experience. One base that also carries compliance, preferred doctor and doctor-written diagnosis — and later can plug in external vendors (MyRx) or consultation-call revenue (*further out*).

Moves Cuts the doctor-call-step discard → **OP2OD**; lifts consult quality → customer experience; removes a big chunk of the incident load → the **foundation** that stops the platform breaking.

More initiatives to come — HA Substitution, High-Intent Reconnection, Incident-to-Order Tracking, plus light mentions of Voicebot Ph2, Discard Intelligence and the CPO cost-view.

HOW Q2 EXECUTES

INITIATIVE	PRIORITY / SIZE	JULY	AUGUST	SEPTEMBER	METRIC
------------	-----------------	------	--------	-----------	--------

Improve Order Completion — OP2OD

Checkout & Payment Journey	P0 / XL	Discovery/Solution	Delivery	Delivery/GTM	SR, Conv Rate, OP2OD
Discard Intelligence (Analytics)	P0 / TBD	Discovery	Delivery	Levers identified	TBD
Personalised HA Substitution	P1 / M	Discovery/Solution	Delivery/GTM	Measure	GM%
Ring AI Cart Recovery	P0 / M	Build	Rollout/Scale	Scale	iOPD, CPO

CPO reduction

CSR Voicebot Phase 2	P0 / M	Build	Delivery/GTM	Measure	CSR CPO
High-Intent Reconnection	P1 / L	Discovery	Delivery/GTM	Measure	OP2OD, Cx NPS, HA TAT
TM Wallet Trust Revamp	P1 / L	Discovery	Delivery	—	Wallet ITO, NPS
Returns & Refunds	P0 / XXL	Discovery	Delivery/GTM	Delivery/GTM	Refund TAT, CSR CPO 10L/month

Foundations

Doctor Experience Platform	P0 / XL	Discovery	Delivery/GTM	Delivery/GTM	Dr CPO (retention)
-----------------------------------	---------	-----------	--------------	--------------	--------------------

INITIATIVE	PRIORITY / SIZE	JULY	AUGUST	SEPTEMBER	METRIC
Incident-to-Order Tracking	P1 / XXL	Discovery	Discovery	Delivery/GTM	ITO
Core Configuration & Controls	P0 / L	Discovery/Solution	Delivery	—	Config lead time · incident rate

Stage placement is a working plan, not a commitment — it firms up as the middle is written and capacity is allocated.

Capacity (working model): 6 delivery engineers — 2 Payment, 4 Portal/Platform · 1 engineering lead (oversight, not delivery) · 1 SBA + 1 BA · 1 SPM · 0 execution PM/APM. Allocation still open.

APPENDIX · FOR ANALYTICS & FINANCE

Making Q2 measurable

Reference for the working session with Analytics & Finance — not CEO-core. Q2 turns measurable only when these data products exist; each one unblocks a decision. OP2OD is already defined; everything below is open. One SBA and one BA support this contract; allocation is unresolved.

WHAT WE BUILD

THE DECISION IT UNLOCKS

P0 — needed first	
CPO cost view — collate every cost centre per order delivered: RTO, COD-DR, HA, CSR, AI cost and the rest, into one CPO picture.	A directional cost map — where cost actually sits, so we know where to attack and what to take into discovery.
Discard diagnostic — attribute discards by stage, payment / return path, category, call cohort, root cause, recoverability, value, owner.	Which discard cohorts earn Q2 action — a ranked lever pipeline.
Payment funnel — entry → payment → return state → delivery, with broken-return attribution.	Checkout priorities and clean online-vs-COD attribution. No baseline exists today.

WHAT WE BUILD

THE DECISION IT UNLOCKS

Contact / friction model — contact → customer → order → reason → channel → resolution → repeat → callback TAT.

CRM scope, FCR and repeat visibility; friction measured backwards from the customer.

Voicebot cost bridge — eligible volume, containment, transfer, FCR, repeats, quality, human effort and cost avoided.

Whether the ₹9/order CSR CPO holds. Scale or iterate Voicebot Phase 2.

Doctor baseline — consultation time, workflow correctness, rework, intervention rate.

Doctor Experience targets and the September readout.

Returns & Refunds funnel — deduplicated cases, repeats, team + cost, eligibility, approval, pickup, refund, TAT.

The refund-automation milestone. Call volume alone can't size unique friction.

P1 — next

Ring AI experiment — eligible pool, treatment / control, conversion, delivery, cost, incrementality.

Any Ring AI impact claim; scale / iterate / stop. (12.5% is a hypothesis, not a result.)

Reconnection cohort — deduplicated signal → reconnect → confirmation → delivered order.

Size recoverable high-intent demand; validate the ~250 repeat-callback signal.

Substitution cohort — acceptance → delivery → repeat / LTV → GM contribution.

Whether personalised substitution creates durable value, not one-time acceptance.

Wallet correctness — creation → sync → eligibility → application → failure, plus linked support contacts.

Wallet priority and its success proof. No quantified result today.

Incident / capacity model — reconciled incident activity and engineering effort on recurring work.

Any capacity-release claim; validates or kills the ~50% bandwidth estimate.

Configuration baseline — change volume, lead time, manual effort, failure rate, rollback.

Sizing Core Configuration & Controls.

Trust measurement — one headline Trust metric, or Wallet correctness + R&R resolution as proofs.

The final Trust treatment. Open decision — we won't invent a trust score.